

Remote Invocation

To do ...


- ❑ Today
- ❑ Request-reply, RPC, RMI
- ❑ Next time: Indirect communication

Beyond message passing

- All IPC in distributed systems is based on low-level message passing
 - As we discussed in the last two lectures
- A bit too low level
 - Modern distributed systems can have 10^3 - 10^6 processes scattered around


Beyond message passing

- Need a higher-level of abstraction – send/receive exposes communication
- Are there programmer-friendlier models?
 - Today ...
 - Request-reply – patterns in msg passing with little support for request/reply interactions
 - Remote procedure calls – extending procedure calls
 - Remote method invocation – ... to remote objects



Request-reply protocols

- We can describe it based on three primitives
 - `doOperation` – used by client to invoke operation
 - Args specify remote server and arguments; result is a byte array
 - After sending it, client issues a receive to get the reply
 - `getRequest` – used by server to get request
 - `sendReply` – used by server to send reply
 - When received by client, original `doOperation` is unblocked



Request-reply protocols

- Normally, synchronous (client blocks) and reliable
 - Asynchronous is also possible
- Synchronous and asynchronous
 - Sender continues (asynchronous) or blocks (synchronous) until request has been accepted
 - Points of synchronization: (1) at request submission, (2) at request delivery or (3) after processing
- Reliability, two concerns
 - Integrity – msg arrives uncorrupted and without duplication
 - Reliability – msg arrives despite some packet drops

Request-reply protocols

- For reliability or request-reply communication
 - Messages need a requestId and a process identifier
- Failures partially depend on transport
 - TCP or UDP
 - Over UDP, omission and out-of-order issues
 - Process may also fail (crash failures)
- To handle omission failures – timers
 - For duplicate messages, msg id (keep a history) or idempotent operations

messageType: int (req/reply)
requestId: int
remoteReference
operationId: int
arguments: byte[]

Request-reply protocols

- Exchange styles – different behavior in front of failures
- *Request (R)* – When client doesn't need confirmation, asynchronous (typically over UDP)
- *Request-reply (RR)* – Useful for most client-server exchanges
 - No need special ack, server reply is an implicit ack
- *Request-reply-acknowledge Reply (RRA)* – Server can clean history

Name	Client	Server	Client
R	Request		
RR	Request	Reply	
RRA	Request	Reply	Ack reply

Using TCP or UDP to implement RR

- Client-server exchange can be built on UDP or TCP (or any other transport, of course)
- To avoid implementing multi-packet protocols, TCP
 - TCP reliability means no need for retransmissions, duplicate filtering or history
 - No problem with large transmissions, flow-control handles it
 - If multiple exchanges, connection overhead applies once
- If you can live without all this, maybe a more efficient protocol over UDP
 - Sun NSF transmits fixed-size blocks between client/server
 - All operations are idempotent, so no need for history

Data representation and marshalling

- Processes keep information in data structures
 - records, arrays, strings, trees ...
- But IPC is in msgs, sequences of bytes
 - TCP/UDP gives the mechanisms to send sequences of bytes
 - Processes need a protocol to make the exchange meaningful
 - To serialise this data into a streamy of bytes to write it, and deserialise it to read it: marshalling and unmarshaling

Data representation and marshalling

- **Marshalling/unmarshalling**
 - Assembling/disassembling process' data for transmission
 - Client and server may have different data representations
 - Both need to properly interpret msg to transform it into machine-dependent representation
 - Agree on encoding
 - How are basic data values represented (integers, floats, ...)
 - How are complex data values represented (arrays, unions)
 - Intermediate language or source's representation
- **Multiple external representation alternatives**
 - Sun's XDR, Corba, XML, ASN.1, Google's protocol buffer, JSON, Gob (Go specific)

An example with ASN.1*

```
package main
```

```
import (
```

```
    "bytes"
```

```
    "encoding/asn1"
```

```
    "net"
```

```
    "os" ...
```

```
)
```

```
func main() {
```

```
    if len(os.Args) != 2 {
```

```
        fmt.Fprintf(os.Stderr, "Usage: %s host:port", os.Args[0])
```

```
        os.Exit(1)
```

```
    }
```

```
    service := os.Args[1]
```

```
    conn, err := net.Dial("tcp", service)
```

```
    result, err := readFully(conn)
```

```
    var newtime time.Time
```

```
    _, err1 := asn1.Unmarshal(result, &newtime)
```

```
    fmt.Println("After marshal/unmarshal: ", newtime.String())
```

```
    os.Exit(0)
```

```
}
```

Client code (missing error checking!)

HTTP as an example

- HyperText Transfer Protocol, on top of TCP
 - Specifies msgs exchanged, formats, methods, arguments and results, representation for marshaling ...
 - Content negotiation – clients state format they can accept
 - Password-style authentication
- A fixed set of methods, some well-known ones
 - GET – Requests resource or run program pointed to by URL
 - HEAD – Same as GET but returns only metadata
 - POST – Provides data, depending on function supported by the program specified by the URL (e.g., posting a msg)
 - PUT – Requests to store data with the given URL as ID
 - Others: DELETE, OPTIONS, TRACE

HTTP as an example

- Clients invoke methods to be applied to resources at the server (given by the URL)
- Msgs marshalled into ASCII text strings
- Connections
 - Client interaction in version 1.0
 - Client request a connection at default (or given) port
 - ... sends request msg to server
 - Server sends reply
 - Connection is closed
 - Setting/closing a connection per request is costly
 - HTTP 1.1 uses persistent connections
 - Can be close by client, server or after being idle for a while

HTTP as an example

```
$ telnet www.golang.org 80
Trying 64.233.191.141...
Connected to golang.org.
Escape character is '^I'.
GET /index.htm HTTP/1.1
host: www.golang.org

HTTP/1.1 302 Found
Location: http://golang.org/index.htm
Date: Thu, 09 Apr 2015 14:21:32 GMT
Content-Type: text/html; charset=UTF-8
Server: Google Frontend
Content-Length: 224
Alternate-Protocol: 80:quic,p=0.5
```

```
<HTML><HEAD><meta http-equiv="content-type"
content="text/html; charset=
<TITLE>302 Moved</TITLE></H
<H1>302 Moved</H1>
The document has moved
<A HREF="http://golang.org/
</BODY></HTML>
...
```

```
package main

import (
    "net/http" ...
)


func main() {
    response, err := http.Get("http://golang.org")
    if err != nil { ...
    defer response.Body.Close()
    contents, err := ioutil.ReadAll(response.Body)
    if err != nil { ...
    fmt.Printf("%s\n", string(contents))
}
```

Using telnet

Using go

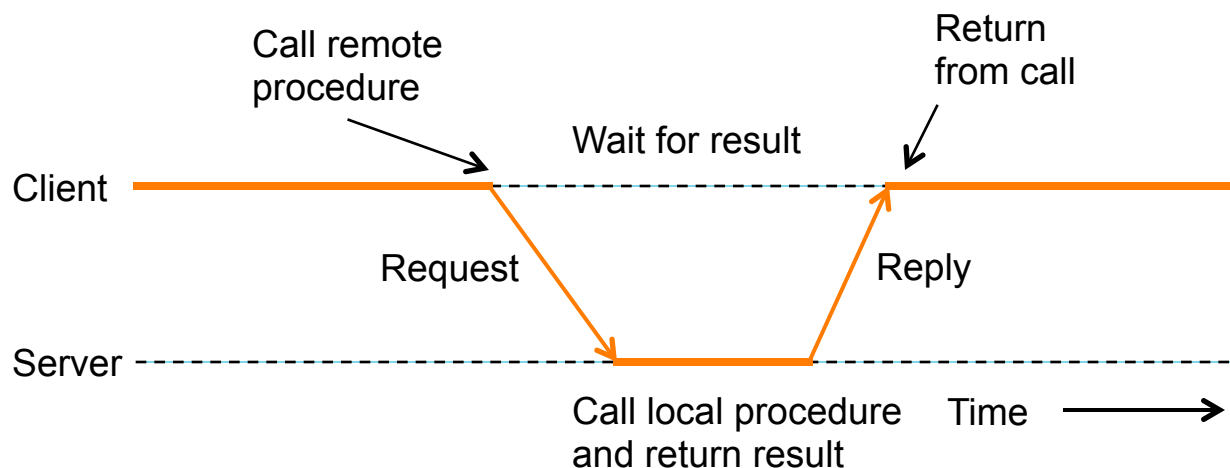
Back in 5'

- ... RPC and RMI



Remote Procedure Call (RPC)

- Earliest and best known example of a more programmer friendly model [Birrell and Nelson '84]
- Some observations
 - Developers are familiar with simple procedure model
 - Well engineered procedures operate in isolation
 - No fundamental reason not to execute procedures on a separate machine
- Can hide sender/receiver comm. using proc calls?



RPC details

- RPC promote programming with interfaces
 - Better abstractions & maintainability, language independence
 - Interface specification with lang independence – Interface Definition Languages (e.g., XDR, Corba IDL)
- RPC, local procedure calls and transparency
 - Parameter passing and global variables
 - Copy in/copy out semantics – while procedure is being executed, nothing can be assumed about parameter values
 - All data to be worked on is passed by parameters; no ref to globals
 - How about pointers?
 - Copy/restore, no call-by-reference
 - Remote reference for more complex structures
 - Failures and latency

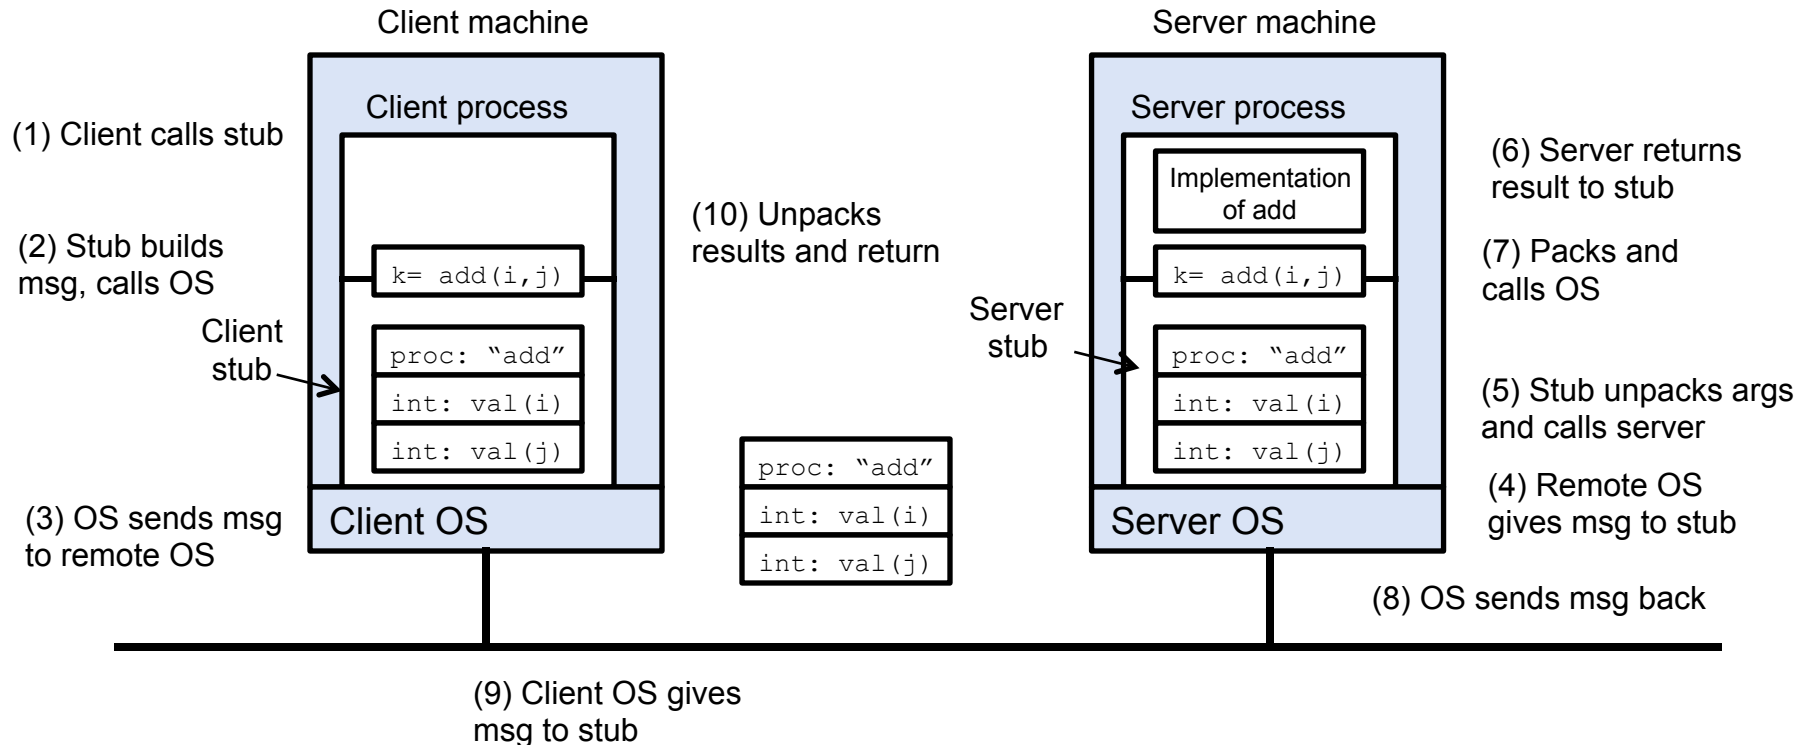
RPC details

- RPC call semantics
 - Depending on fault tolerance measures:
 - Retransmit request until getting a reply or decide server failed
 - Duplicate filtering at the server
 - Re-execute procedure or retransmit reply, keeping history of results at the server

Call semantics	Fault tolerance measures		
	Retransmit request	Duplicate filtering	Re-execute/retransmit
Maybe	No	NA	NA
At-least-once	Yes	No	Re-exec
At-most-once	Yes	Yes	Retransmit reply

- *What are the semantics of local procedure calls?*

Basic RPC implementation and operation

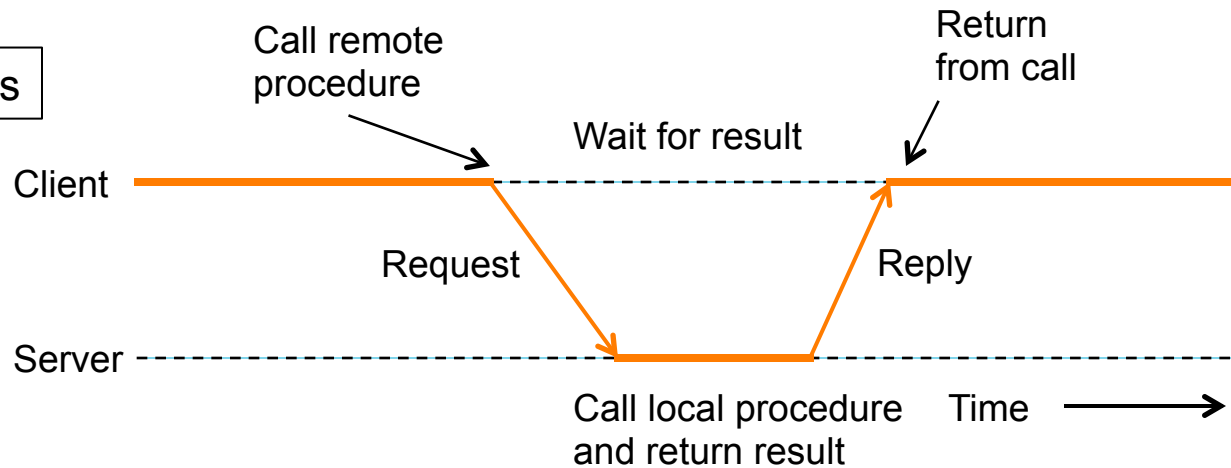


RPC details

- Runtime is given
 - RPCRuntime was part of Cedar in the original RPC system
- Programmers writes client and server
- Client and server-stub are user generated
 - Based on the interface specification
 - By *Lupine* in the original
- A binder for clients to find where to connect
 - Binder runs on a well-known-port
 - Manage table of references/ports for each service

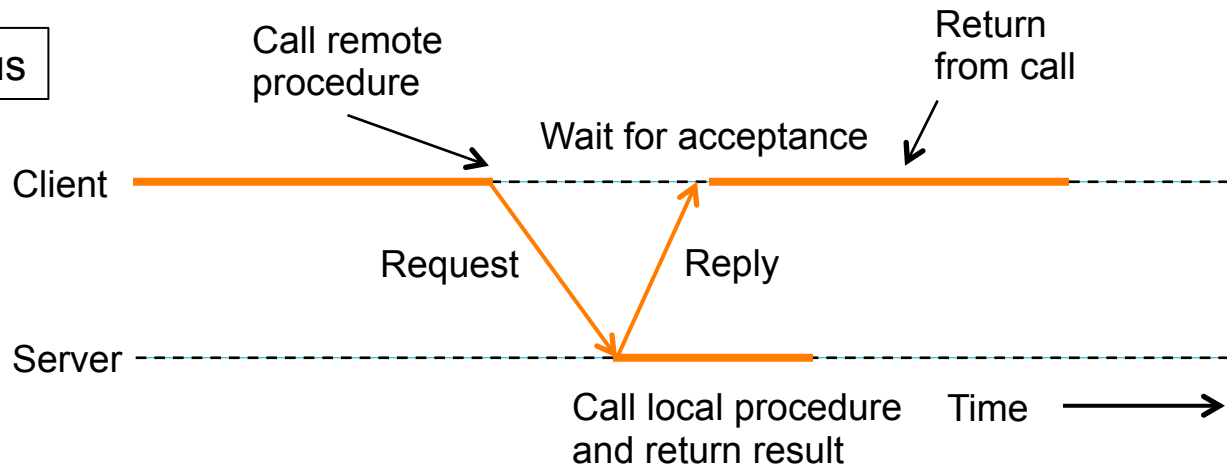
Asynchronous RPCs

Synchronous



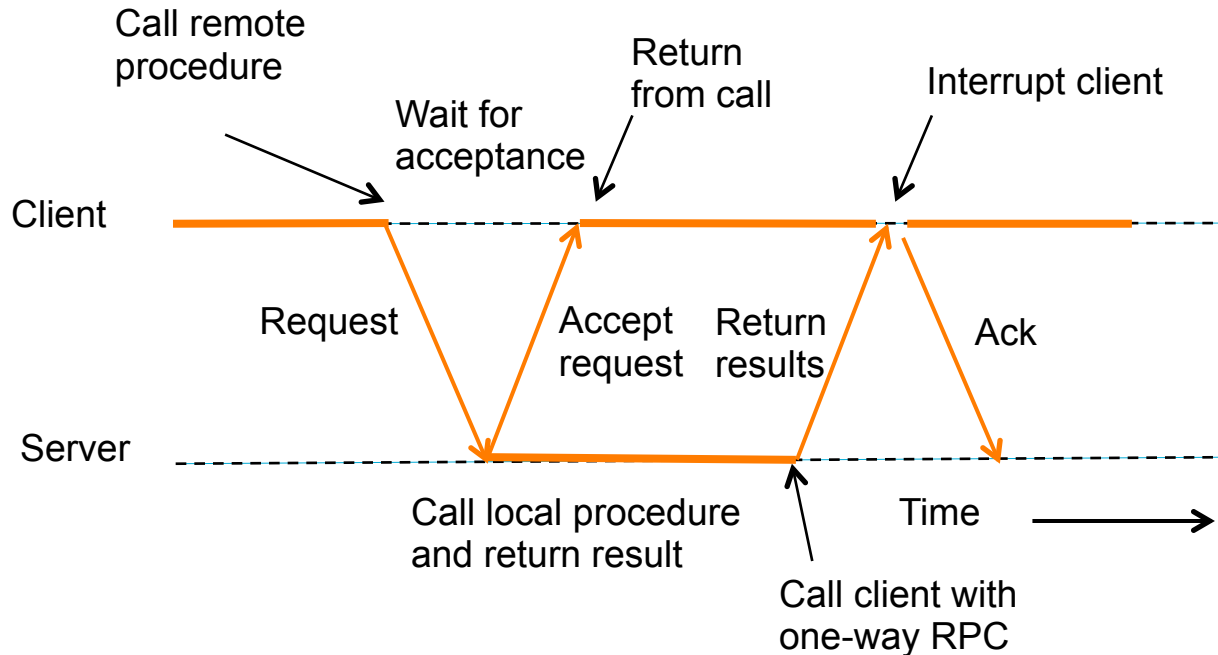
- Get rid of the strict request-reply behavior, but let the client continue w/o waiting for server's answer

Asynchronous



Deferred synchronous RPCs

- Combining two asynchronous RPC is sometimes also referred to as deferred synchronous RPC



- A variation – Client can also do a (non)blocking poll at the server to see whether results are available

Sun RPC

- Defined in RFC 1831
- Designed for client-server communication in the Sun Network File System (NFS)
- Run over UDP or TCP, using at-least-once semantics
- Sun XDR as external data representation and IDL
- Interfaces are identified by program and version number
- Binder (port mapper) for clients to find where to connect
- Authentication through fields in the request/reply msgs

Remote Method Invocation

- RMI extends RPC into the world of distributed objects
 - As RPC, programming with interfaces
 - ... also built on top of request-reply, offers similar call semantics
 - ... and similar level of transparency
- But
 - Programmer can use OO programming features (objects, classes, inheritance ...)
 - All objects have an object reference; refs can be passed on as parameters (first class values)
 - Not just parameter passing by value, good for complex parameters
 - For distributed objects, remote object refs and remote interfaces

Remote Method Invocation

- With OO, state partition among processes as objects
- If using a client-server model,
 - Objects managed by servers, invoked by clients through RMI
 - Objects could also be replicated and/or migrated for reliability, availability or performance
- Implementing RMI
 - Similar to RPC, a proxy object, two communication modules and a dispatcher & skeleton
- With distributed objects, distributed garbage collection

Distributed garbage collection

- One way to implement it – cooperating local collectors
 - Server keeps list of processes holding remote refs to its objects
 - When a client first receive a remote ref. to a remote object, adds itself as holder at server (*extra invocation*) and creates a proxy
 - Server adds clients to holders
 - When client garbage collects proxies for remote object, removes itself from holders at server (*extra invocation*) then deletes proxy
 - Server removes client from holders
 - Java's approach
- Keeping resources at servers and leases
 - What to do if clients go away? Set up leases granting the use of resources for a fixed period of time

Summary

- Powerful primitives can makes (distributed) programming them a lot easier
- Procedure calls
 - Simple way to pass control and data
 - Elegant and transparent way to distribute applications
 - Not the only way
- Hard to provide true transparency
 - Failures, performance, memory access, ...
- How to deal with hard problems – let the programmer do it – “worse is better” (Richard Gabriel’s)