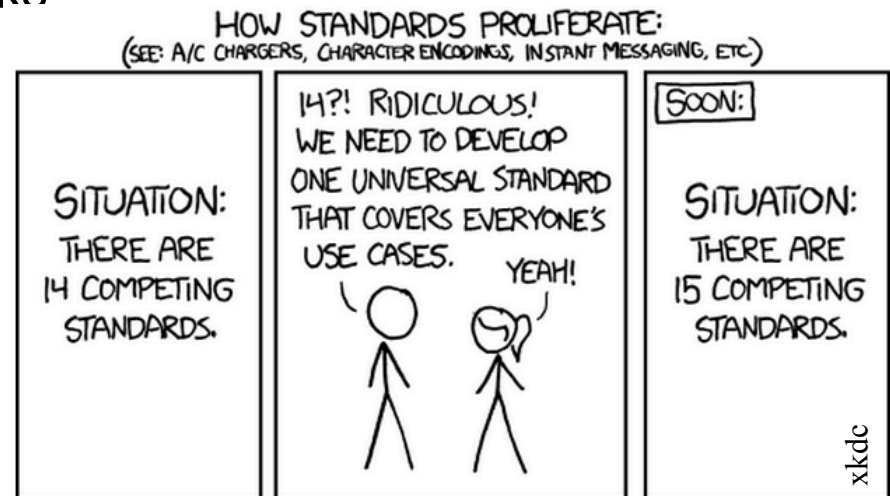


Indirect Communication

To do ...

- ❑ Today
- ❑ Space and time (un)coupling
- ❑ Common techniques
- ❑ Next time: Overlay networks



Direct coupling communication

- With R-R, RPC, RMI
 - Space coupled – Sender knows the identity of the receiver and vice-versa
 - If server fails, hard to replace; clients must explicitly deal with that
 - Time coupled – Server and receiver must both exist at the time of communication

	Time-coupled
Space coupling	Communication directed to a given receiver(s) that must be available at the time <i>e.g. Messaging passing, RPC</i>

Indirect communication

- Indirect communication
 - Through an intermediary with no direct coupling between senders/receivers
 - Different types, with differences in nature of the intermediary
 - ... and type of coupling
- Forms of (un)coupling
 - Space – (No) need to know the identity of the other
 - Time – (No) need to exist at the same time

Space and time (un)coupling

- Space uncoupling
 - No need to know the identity of the other party
 - Can change, update, replicate, move senders/receivers
 - E.g., IP multicast
- Time uncoupling
 - No need to exist at the same time
 - It's ok if either party gets disconnected for a bit
 - E.g., Mailbox
 - Not the same as asynchronous – *Why not?*

Space and time un/coupling

	Time-coupled	Time-uncoupled
Space coupling	<p>Communication directed to a given receiver(s) that must be available at the time</p> <p><i>e.g. Messaging passing, RPC</i></p>	<p>Sender(s) and receiver(s) can have independent lifetimes</p> <p><i>e.g. Mailbox</i></p>
Space uncoupling	<p>Sender does not need to know ID of receiver but they must exist at the same time</p> <p><i>e.g. IP multicast</i></p>	<p>Sender does not need to know ID of receiver; sender(s) and receiver(s) can have independent lifetimes</p> <p><i>e.g. Message oriented</i></p>

Examples of indirect communication


- Group communication
 - An abstraction over multicast communication
 - Space uncoupled
- Publish-subscribe systems
 - The most widely used indirect comm. techniques
 - Space uncoupled and possible time uncoupled
- Message queues
 - Space and time uncoupling through a msg queue
- Shared memory
 - Distributed shard memory and tuple spaces

Group communication

- Sender communicates with a group, as a whole, without knowing the identity of members
 - An abstraction over multicast (IP or overlay)
- Group comm. typically to process groups
 - Some work on object groups – a collection of objects, typically instances of the same class
- Some common uses
 - Reliable dissemination (e.g., financial reports)
 - Collaborative applications (e.g., multiuser games)
 - Highly-available services
 - System monitoring and management

Group communication

- Groups of processes
 - Processes can join/leave
 - A message to the group reaches all (broadcast)
- Not just programmer convenience



More efficient use of
bandwidth – just once
per network link

Groups and group management

- Process and object groups
 - Messages, typically unstructured byte arrays, are delivered to processes ~ socket
 - A collection of objects that process the same set of invocations concurrently; client invokes a method once on a local proxy
- Groups may be
 - Closed or open – only members or anyone can send to group
 - Overlapping or not – processes can be members of 1+ group
 - Synchronous or asynchronous
- Group membership
 - Membership service provides interface for membership changes, failure detection and notification

Reliable and ordered multicast

- Reliable


- From one-to-one communication
 - Integrity – msg received is the one sent and no msg is delivered twice
 - Validity – any outgoing msg is eventually delivered
- Agreement – if msg is delivered to one, it is delivered to all

- Ordered

- FIFO ordering – source ordering, preserve order of the sender
- Causal – causally related msgs arrive in the same order everywhere; if a msg *happens before* another msg this so called *causal relationship* is preserved in the delivery
- Total ordering – All msgs arrive in the same order everywhere

JGroups toolkit as an example

- An example based on Birman and van Renesse' work on ISIS, Horus and Ensemble
- Includes channels (handle onto a group), building blocks and a composable stack
 - Every module can be stack over/bellow any other
 - Not all stacks make sense, of course
 - CAUSAL – causal ordering
 - FRAG – configurable packetization
 - GMS – group membership system to maintain consistent view of the group
 - MERGE – Network partitions and group merges
 - ...

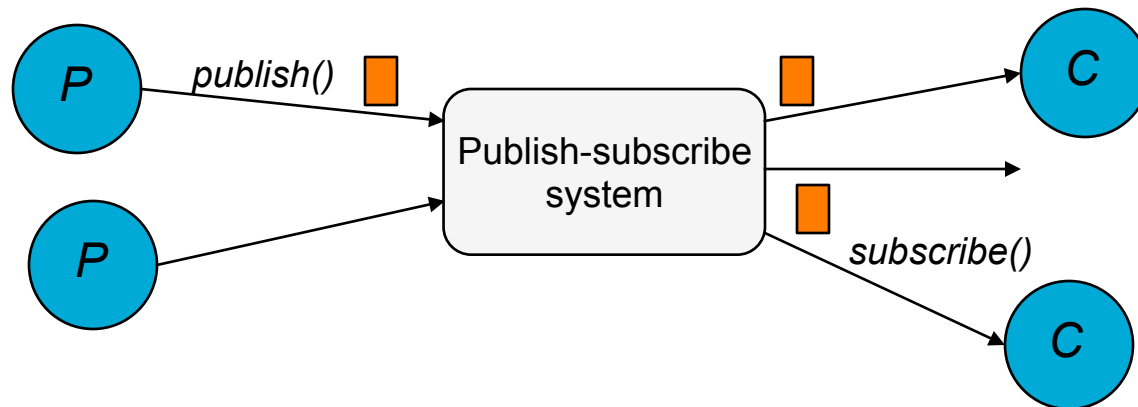


Examples of indirect communication

- Group communication
 - An abstraction over multicast communication
 - Space uncoupled
- Publish-subscribe systems
 - The most widely used indirect comm. techniques
 - Space uncoupled and possible time uncoupled
- Message queues
 - Space and time uncoupling through a msg queue
- Shared memory
 - Distributed shard memory and tuple spaces

Publish-subscribe

- AKA distributed event-based systems
 - The most widely used of all indirect communication models
 - E.g., CORBA Events, TIB Rendezvous, Scribe, Echo, ...
- Publishers and subscribers
 - Publishers publish events, subscribers subscribe to them
- The pub/sub system job
 - Match subscriptions with published events, ensure delivery



- Example applications
 - financial systems, live feeds, monitoring apps, ...


Programming and subscription models

- Simple programming model
 - Publishers *publish(e)*, subscribes *subscribe(f)* where *f* is a filter on the type of events they care for, *unsubscribe()*
- Different subscription models
 - Channel-based
 - Basic, publishing to named channels
 - Topic or subject based
 - Notifications are expressed in terms of a number of fields; one field denotes the topic
 - Content-based
 - Allows subscription over a range of fields
 - Other types explored
 - Type-, context- and concept-based and more complex event processing

Publish-subscribe – implementation

- Goal – efficient delivery of the right events to the right subscribers with appropriate security considerations
- Some design options

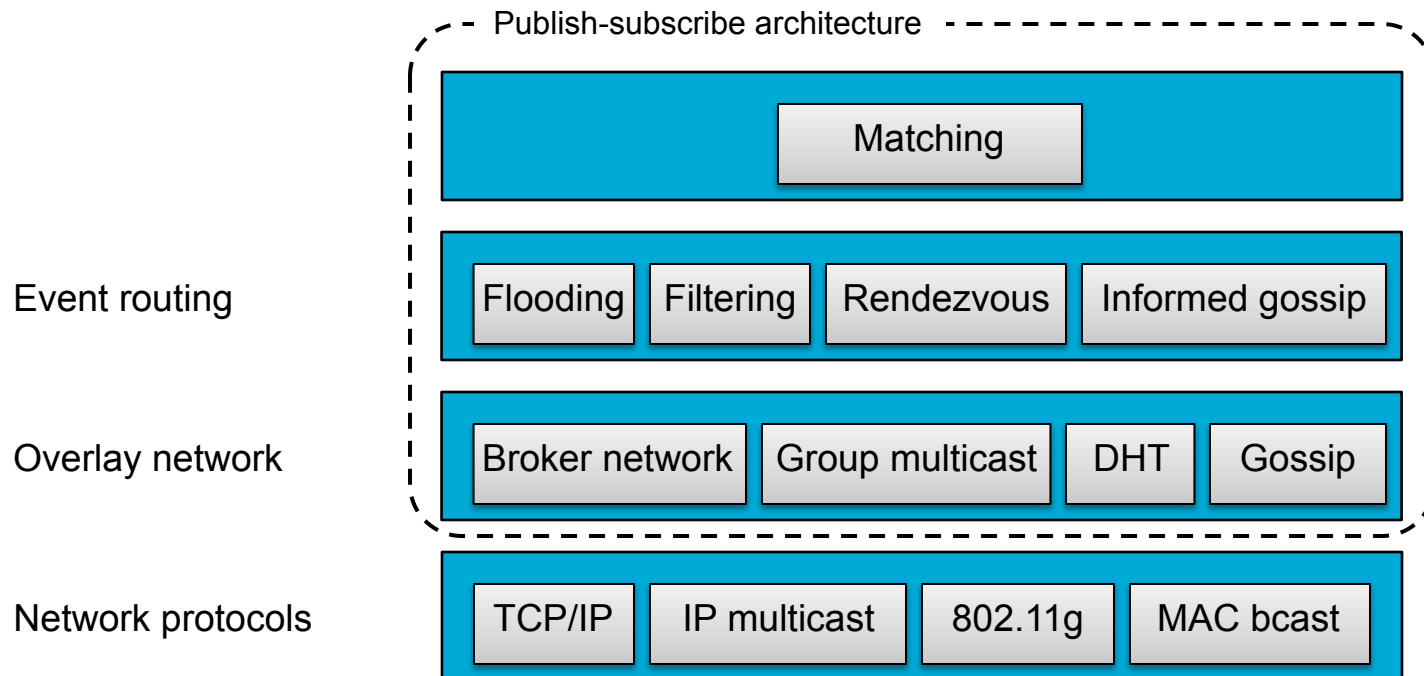
- Centralized/distributed
 - Centralized event broker
 - Network of brokers



- Full P2P – not distinction between publishers and subscribers, i.e., everyone is a broker
- Routing options ...


Architecture of publish-subscribe systems

- Routing options
 - Flooding – send to all, matching done at the subscriber
 - Filtering – every node in the network of brokers does filtering-based routing
 - Rendezvous – a node responsible for matching notifications and subscribers




Back in 5'

- ... message queues and shared memory



(Distributed) message queues

- Point-to-point comm. through an intermediary queue
 - Senders place msgs into a queue, receivers removed them
 - Queues correspond to buffers at communication servers
 - E.g., IBM WebSphere MQ, Java Messaging Service, Oracle's Stream Advanced Queuing




(Distributed) message queues

- Details on messages
 - Typically include dest queue, priority, delivery mode, and body
 - In Oracle's AQ, messages are rows in a DB table/queue
 - Messages are persistent
 - Typical queuing policies FIFO and priority-based
- Use for app integration – broker takes care of application heterogeneity
 - Transforms incoming messages to target format
 - Often acts as an application gateway
 - May provide subject-based routing capabilities

(Distributed) message queues

- Styles of receive
 - Blocking – Block until an appropriate message is available
 - Non-blocking – Polling to see if a message is available
 - Notify – Notify when a message arrives
- Centralized and distributed message queues
 - In WebSphere MQ, queues are managed by *queue managers*
 - Queue managers can be inter-connected as brokers in pub-sub



Examples of indirect communication


- Group communication
 - An abstraction over multicast communication
 - Space uncoupled
- Publish-subscribe systems
 - The most widely used indirect comm. techniques
 - Space uncoupled and possible time uncoupled
- Message queues
 - Space and time uncoupling through a msg queue
- Shared memory
 - Distributed shard memory and tuple spaces

Shared memory approaches

- Distributed share memory
 - Allow networked computers to share memory
 - How to make distributed memory appear local?
 - Leverage MMU
 - Page fault handler invokes DSM protocol
 - Bring page from a remote node instead of from HD
 - Of course, underneath it all – message passing
- Compare with message passing
 - No need to marshalled/unmarshalled data, send/receive, ...
 - Synchronization via typical shared-memory programming constructs like locks
 - Potentially easier to program for with a performance cost

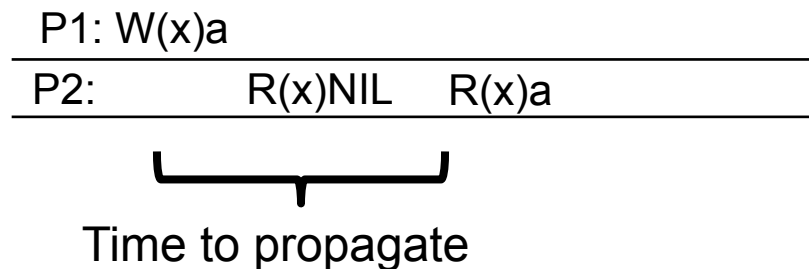
Shared memory approaches

- Simplest design
 - Each virtual page in one machine at a time (no caching)
 - A directory keeps track of things, potentially a bottleneck
 - Distributed directory – hash(page#)
 - Design issues
 - Size of the page
 - Caching and consistency models





Shared memory and consistency

- Consistency model
 - When modifications to data may be seen at a given processor
 - Defines the programmer's view, placing restrictions on what values can be returned by a read (a contract)
 - Determines what optimizations are possible
- E.g., sequential consistency
 - Some basic notation
 - $W_i(x)a$ – process P_i wrote value a to x
 - $R_i(x)b$ – process P_i read value b from x



Sequential consistency

- *Result of execution – as if*
 - *operations of all processes were executed in some sequential order, and ...*
 - *the operations of each process appear in this sequence in the order specified by its program*
 - i.e., Any valid interleaving of ops is OK, but all processes see the same interleaving




- **Linearizable** – interleaving is consistent with real time at which operations occurred in the actual execution

The burden of sequential consistency

- Processor must ensure that previous memory operation is complete before proceeding to the next
- So ...
 - Determine completion of write; get ack for all
 - If caching, write invalidates or updates all cached copies
 - Hold off on read requests until all writes are complete
- Maybe we can relax this a bit if next steps don't depend on the value
 - Causal consistency ...

Shared memory – tuple spaces

- First introduced with *Linda* by D. Gelernter
 - Adopted by IBM Tspaces, JavaSpaces, etc.
- Programming model
 - Processes communicate through a tuple space, a shared collection of tuples
 - Tuple – a sequence of 1+ typed data fields
 - Operations
 - Write – adds a tuple
 - Read – returns the value of a tuple w/o changing the tuple space
 - Take – returns the value of a tuple and removes the tuple
 - For read/take, give a template; system returns a tuple that matches
 - Tuples are immutable - to modify a tuple, take it and write a new one



Shared memory – tuple spaces

- Original Linda model had a single, global tuple space
 - Not optimal for a large system, e.g., aliasing of tuples
 - Aliasing – read/take matching tuples by accident
 - Following systems use multiple tuple spaces and some allow the dynamic creation of tuple spaces
- Linda was anticipated as a centralized system
 - Performance and reliability concerns
 - Following systems support distributed tuple spaces
 - Different approaches from state machine replication and tuple-specific approaches to simple partitioning of the tuple space

Summary

- The power of indirection in communication – communication through an intermediary
 - Uncoupling in space and/or time

	Group	Pub/sub	MQ	DSM	Tuples
Space uncoupled	Yes	Yes	Yes	Yes	Yes
Time uncoupled	Possible	Possible	Yes	Yes	Yes
Style	Comm	Comm	Comm	State	State
Comm pattern	1-m	1-m	1-1	1-m	1-1/1-m
Scalability	Limited	Possible	Possible	Limited	Limited
Associative	No	Content-based only	No	No	Yes